



Data Science II Assignment Individual

WS 2025

Bachelor: Mechatronik, Design & Innovation

5th Semester

Lecturer: Daniel T. McGuiness, Ph.D

Groupe: BA-MECH-23-MLDS

Author: Lukas Heiss

January 20, 2026

Contents

1	Introduction	1
2	Q1 - Visualising Olivetti Faces	2
2.1	Code	2
2.2	Results	3
3	Q2 - Clustering the Faces	4
3.1	Code	4
3.2	Results	6
4	Q3 - Gaussian Mixture Models	7
4.1	Code	7
4.2	Results	9
5	Conclusion	11
List of Figures		III
List of Tables		IV
References		V

1 Introduction

This assignment addresses the processing of the Olivetti Faces dataset using a combination of unsupervised and supervised learning methods. Starting with a basic K-Means clustering for structure analysis in section Q1 2, a further optimized pipeline with decision trees is developed, whose ideal parameters are determined through a systematic silhouette analysis in section Q2 3. The conclusion involves the use of principal component analysis (PCA) for efficiency improvement, as well as the application of Gaussian Mixture Models (GMM) to evaluate the generative properties of the model in creating new faces and detecting image anomalies in section Q3 4.

2 Q1 - Visualising Olivetti Faces

The first part of this project, the Olivetti Faces dataset is loaded and examined. The goal is to apply the K-Means algorithm to divide the 400 images into 40 clusters based on their pixel data, which corresponds to the number of 40 people contained in the dataset. The analysis focuses on how well the algorithm recognizes similarities in facial features. Finally, the distribution of images across the individual clusters is evaluated in order to assess and understand the grouping of the clusters.

2.1 CODE

Listing 1: Q1 - Olivetti Faces Clustering with K-Means

```
1 from sklearn.datasets import fetch_olivetti_faces
2 from sklearn.cluster import KMeans
3 from sklearn.discriminant_analysis import StandardScaler
4 from sklearn.model_selection import train_test_split
5 import matplotlib.pyplot as plt
6 import numpy as np
7 from sklearn.tree import DecisionTreeClassifier
8 from sklearn.metrics import accuracy_score
9
10 # ++++++ Q1 - Olivetti Faces Clustering with K-Means
11 # ++++++
12 """
13 We load the Olivetti faces dataset, and use K-Means for clustering. We
14 also visualize the faces within a
15 selected cluster. Then, we count and print how many images are in each
16 cluster.
17 """
18
19 # --- loads Olivetti DS ---
20 data_faces = fetch_olivetti_faces()
21 X = data_faces.data
22 y = data_faces.target
23
24 # --- Split dataset in training and test set ---
25 # stratify mixes the classes in both sets equally, so that all persons
26 # are represented in both sets
27 X_train, X_test, y_train, y_test = train_test_split(
28     X,
29     stratify=y,
30     test_size=0.2,
31     random_state=42
32 )
33
34 #--- Model with k-means clustering ---
35 # 400 images of 40 different persons = 40 clusters
36 kmeans = KMeans(n_clusters=40, random_state=42)
37 kmeans.fit(X_train)
38
39 y_pred = kmeans.predict(X_test)
40 accuracy_Kmeans = accuracy_score(y_test, y_pred)
41 print(f"Kmeans accuracy: {accuracy_Kmeans:.2f}")
42
43 # --- Function to plot the faces within a cluster ---
44 var = 0 # Cluster number to visualize
45 cluster_labels = kmeans.labels_
46 clusters = X_train[cluster_labels == var]
```

```

46 # --- Counts how many images are in each cluster ---
47 cluster_counts = np.bincount(kmeans.labels_)
48 for i, count in enumerate(cluster_counts):
49     print(f"Cluster {i} has {count} Images")
50
51 # --- Ploting the pictures in cluster var ---
52 plt.figure(figsize=(10, 2))
53 for i in range(min(20, len(clusters))):
54     plt.subplot(1, 20, i + 1)
55     plt.imshow(clusters[i].reshape(64, 64), cmap='gray')
56     plt.axis('off')
57 plt.show()

```

2.2 RESULTS

The K-Means algorithm achieved an accuracy of 0.0625. The distribution of images across the 40 identified clusters is summarized in Table 2.1 and the pictures of the faces in the first cluster (Cluster 0) are shown in the Image 2.1.

Table 2.1: Distribution of Images per Cluster

Cluster	Count	Cluster	Count	Cluster	Count
Cluster 0	14	Cluster 14	17	Cluster 28	9
Cluster 1	11	Cluster 15	4	Cluster 29	9
Cluster 2	14	Cluster 16	3	Cluster 30	2
Cluster 3	14	Cluster 17	4	Cluster 31	11
Cluster 4	6	Cluster 18	4	Cluster 32	4
Cluster 5	8	Cluster 19	8	Cluster 33	8
Cluster 6	11	Cluster 20	5	Cluster 34	17
Cluster 7	5	Cluster 21	13	Cluster 35	5
Cluster 8	10	Cluster 22	6	Cluster 36	6
Cluster 9	8	Cluster 23	5	Cluster 37	5
Cluster 10	16	Cluster 24	12	Cluster 38	7
Cluster 11	4	Cluster 25	2	Cluster 39	2
Cluster 12	7	Cluster 26	6		
Cluster 13	8	Cluster 27	10		



Figure 2.1: Plot of the Faces in the first Cluster.

3 Q2 - Clustering the Faces

In the second section, the focus is on supervised learning to improve the classification of the Olivetti Faces. First, a Decision Tree Classifier is trained directly on the image data to measure accuracy in person recognition. Then, a pipeline is implemented that combines K-Means clustering as a preprocessing step with the Decision Tree. The core of this investigation lies in the systematic silhouette analysis, in which various cluster sizes ($k = 20$ to 160) were tested to determine the optimal size. The implementation is based on the methodological documentation of Scikit-Learn [1], whereby the use of the silhouettescore enabled an objective assessment of cluster quality (see table 3.1, which was also illustrated in the plot 3.1).

3.1 CODE

Listing 2: Q1 - Olivetti Faces Clustering with K-Means

```
1 from sklearn.datasets import fetch_olivetti_faces
2 from sklearn.cluster import KMeans
3 from sklearn.discriminant_analysis import StandardScaler
4 from sklearn.model_selection import train_test_split
5 import matplotlib.pyplot as plt
6 import numpy as np
7 from sklearn.tree import DecisionTreeClassifier
8 from sklearn.metrics import accuracy_score
9 from sklearn.pipeline import make_pipeline
10
11
12 # --- loads Olivetti DS ---
13 data_faces = fetch_olivetti_faces()
14 X = data_faces.data
15 y = data_faces.target
16
17 # --- Split dataset in training and test set ---
18 # stratify mixes the classes in both sets equally, so that all persons
19 # are represented in both sets
20 X_train, X_test, y_train, y_test = train_test_split(
21     X,
22     stratify=y,
23     test_size=0.2,
24     random_state=42
25 )
26
27 #--- Model with k-means clustering ---
28 # 400 images of 40 different persons = 40 clusters
29 kmeans = KMeans(n_clusters=40, random_state=42)
30 kmeans.fit(X_train)
31
32 y_pred = kmeans.predict(X_test)
33 accuracy_Kmeans = accuracy_score(y_test, y_pred)
34 print(f"Kmeans accuracy: {accuracy_Kmeans:.2f}")
35
36
37 # ++++++ Q2 - Olivetti Faces using DT
38 # ++++++
39 """
40 Training a Decision Tree Classifier on the Olivetti faces dataset to
41 classify images of different persons.
42 Select the best cluster size for K-Means using silhouette score analysis.
43 """
44
```

```

43 # --- DT Classifier ---
44 # max_leaf_nodes = 40 because, 40 diff. people
45 dt_Classifier = DecisionTreeClassifier(random_state=42, max_leaf_nodes=
46     40)
47
48 # --- Train the model ---
49 dt_Classifier.fit(X_train, y_train)
50
51 # --- Predict on test set ---
52 y_pred_dt = dt_Classifier.predict(X_test)
53
54 # --- Calculate accuracy ---
55 accuracy_dt = accuracy_score(y_test, y_pred_dt)
56 print(f"DT accuracy: {accuracy_dt:.4f}")
57
58 # --- Decision Tree Classifier Pipeline ---
59 DTKmeans_Pipeline = make_pipeline(
60     KMeans(n_clusters=40, random_state=42),
61     DecisionTreeClassifier(random_state=42)
62 )
63
64 # --- Train the pipeline ---
65 DTKmeans_Pipeline.fit(X_train, y_train)
66
67 # --- Predict Pipeline ---
68 y_pred_DTKmeans = DTKmeans_Pipeline.predict(X_test)
69 accuracy_DTKmeans = accuracy_score(y_test, y_pred_DTKmeans)
70 print(f"DT Kmeans Pipeline accuracy: {accuracy_DTKmeans:.4f}")
71
72 from sklearn.metrics import silhouette_score
73
74 # --- Test via silhouette score to find the best cluster size for this DS
75 # for more Infos: https://scikit-learn.org/stable/auto\_examples/cluster/plot\_kmeans\_silhouette\_analysis.html#sphx-glr-auto-examples-cluster-plot-kmeans-silhouette-analysis-py
76 k_range = range(20, 161, 10) # Test 20, 30, ..., 160 clusters sizes
77 silhouette_scores = []
78
79 for k in k_range:
80     kmeans_test = KMeans(n_clusters=k, random_state=42, n_init=10)
81     labels = kmeans_test.fit_predict(X_train)
82
83     # Callculate the score
84     score = silhouette_score(X_train, labels)
85     silhouette_scores.append(score)
86     print(f"The Silhouette Score of k={k}: {score:.4f}")
87
88 # --- Plot thee silhouett scores ---
89 plt.figure(figsize=(8, 4))
90 plt.plot(k_range, silhouette_scores, "bo-")
91 plt.xlabel("Number of Clusters (k)")
92 plt.ylabel("Silhouette Score")
93 plt.title("Determine the Optimal k")
94 plt.grid(True)
95 plt.show()

```

3.2 RESULTS

The results of the silhouette analysis, shown in Figure 2.1 and detailed in Table 3.1, identify $k = 130$ as the optimal number of clusters for the Olivetti face dataset. At this point, the silhouette score reaches its maximum of 0.2148, indicating the most effective separation of the different facial groups. This achieves an accuracy in the pipeline with a decision tree classifier of 0.625. The continuous increase in the silhouette plot shows that significantly more clusters than the original 40 individuals are required to better capture variations in lighting and head position, providing the classifier with more precise features for more accurate identification.

Table 3.1: Silhouette Scores for different Cluster Sizes (k)

Number of Clusters (k)	Silhouette Score	Number of Clusters (k)	Silhouette Score
$k = 20$	0.1144	$k = 100$	0.2046
$k = 30$	0.1255	$k = 110$	0.2023
$k = 40$	0.1422	$k = 120$	0.2007
$k = 50$	0.1558	$k = 130$	0.2148
$k = 60$	0.1691	$k = 140$	0.1971
$k = 70$	0.1845	$k = 150$	0.1930
$k = 80$	0.1895	$k = 160$	0.1825
$k = 90$	0.1965		




Figure 3.1: Plot of the Silhouette Score test.

4 Q3 - Gaussian Mixture Models

In the third and final part of the task, the focus is on improving the efficiency of the model as well as applying generative methods. Since working with high-dimensional image data (pixel values) is extremely demanding, principal component analysis (PCA) is used initially. The goal is to reduce the data in such a way that 95% of the variance is retained without losing significant information. On this reduced dataset, a Gaussian Mixture Model (GMM) is then implemented. Unlike K-means clustering, the GMM functions as a generative model, allowing the creation of new artificial faces, from the learned distribution. Another focus is on anomaly detection, which examines whether the model is capable of identifying manipulated images, such as an upside-down face, as anomalies based on atypical probability values.

The implementation of the Gaussian Mixture Model as well as the subsequent anomaly detection is based on the approaches of the notebook Hands-on Machine Learning with Scikit-Learn, Keras and TensorFlow [2] and also of the lecture book Data Science II [3]. In particular, the strategy of identifying anomalies through a threshold definition of the probability density was adopted in order to mathematically separate these anomaly cases from the normal ones.

4.1 CODE

Listing 3: Q1 - Olivetti Faces Clustering with K-Means

```
1  from sklearn.decomposition import PCA
2  from sklearn.mixture import GaussianMixture
3  from sklearn.datasets import fetch_olivetti_faces
4  from sklearn.model_selection import train_test_split
5  import matplotlib.pyplot as plt
6  from sklearn.metrics import accuracy_score
7  from sklearn.decomposition import PCA
8  from sklearn.mixture import GaussianMixture
9  import numpy as np
10
11
12 # ++++++ Q3 - Olivetti Faces Clustering with GMM
13 # ++++++
14 """
15 We load the Olivetti faces dataset, apply PCA for dimensionality
16 reduction,
17 and then use Gaussian Mixture Models (GMM) for clustering. PCA is used to
18 reduce dimensionality
19 and therefore improve the time. We also generate new faces and perform
20 anomaly detection by comparing
21 scores of normal and altered faces.
22 """
23
24
25
26
27 # --- loads Olivetti DS ---
28 data_faces = fetch_olivetti_faces()
29 X = data_faces.data
30 y = data_faces.target
31
32
33 # --- Split dataset in training and test set ---
34 # stratify mixes the classes in both sets equally, so that all persons
35 # are represented in both sets
36 X_train, X_test, y_train, y_test = train_test_split(
37     X, y,
38     stratify=y,
39     test_size=0.2,
40     random_state=42
```

```

34 )
35 """
36 PCA (Principal Component Analysis)
37 -----
38 Linear dimensionality reduction using Singular Value Decomposition of the
39 data to project it to a lower dimensional space.
40
41 Parameters
42 -----
43 n_components : float, we keep 95% of variance.
44 random_state : int = 42, 42 is the seed used by the random number
45         generator.
46
47 """
48 # --- PCA for dimensionality reduction ---
49 pca = PCA(n_components=0.95)
50 X_train_pca = pca.fit_transform(X_train)
51 X_test_pca = pca.transform(X_test)
52
53 """
54 GaussianMixture
55 -----
56 Gaussian Mixture Model for clustering.
57 Parameters
58 -----
59 n_components : int = 40, Number of mixture components.
60 random_state : int = 42, 42 is the seed used by the random number
61         generator
62
63 GaussianModel = GaussianMixture(n_components=40, random_state=42)
64
65 GaussianModel.fit(X_train_pca)
66 y_predict = GaussianModel.predict(X_test_pca)
67 accuracy_GMM = accuracy_score(y_test, y_predict)
68 print(f"GMM accuracy: {accuracy_GMM:.2f}")
69
70 # --- Generate new faces ---
71 n_gen = 10
72 X_gen_pca, y_gen = GaussianModel.sample(n_samples=n_gen)
73 X_gen_original = pca.inverse_transform(X_gen_pca)
74
75 # Generated faces and plot the generated faces
76 plt.figure(figsize=(12, 3))
77 for i in range(n_gen):
78     plt.subplot(1, n_gen, i + 1)
79     plt.imshow(X_gen_original[i].reshape(64, 64), cmap='gray')
80     plt.axis('off')
81     plt.title(f"Gen {i+1}")
82 plt.suptitle("Generated Faces")
83 plt.show()
84
85 # Density of the faces
86 densities = GaussianModel.score_samples(X_test_pca)
87
88 # Determine threshold for anomaly detection
89 density_threshold = np.percentile(densities, 4)
90 # select a face to test
91 normal_face = X_test[0]
92 # Rotate the face
93 anomalous_face = np.fliplr(normal_face.reshape(64, 64)).reshape(-1)
94 # Transform both faces using PCA

```

```

95 test_faces_pca = pca.transform([normal_face, anomalous_face])
96
97 # Score both faces
98 test_scores = GaussianModel.score_samples(test_faces_pca)
99
100 # Print results
101 print(f"Threshold: {density_threshold:.2f}")
102 print(f"Normal face score: {test_scores[0]:.2f} -> Anomaly? {test_scores
103 [0] < density_threshold}")
104 print(f"Anomalous face score: {test_scores[1]:.2f} -> Anomaly? {
105     test_scores[1] < density_threshold}")
106
107 # Visualization
108 plt.figure(figsize=(8, 4))
109 plt.subplot(1, 2, 1)
110 plt.title(f"Normal\nScore: {test_scores[0]:.2f}")
111 plt.imshow(normal_face.reshape(64, 64), cmap='gray')
112 plt.axis('off')
113
114 plt.subplot(1, 2, 2)
115 plt.title(f"Anomalous (Flipped)\nScore: {test_scores[1]:.2f}")
116 plt.imshow(anomalous_face.reshape(64, 64), cmap='gray')
117 plt.axis('off')
118
119 plt.show()

```

4.2 RESULTS

The results of the Gaussian Mixture Model show two very different aspects of model performance. The measured accuracy of 0.0125 is predictably low, as the GMM, like K-Means, is an unsupervised learning method. The model clusters faces based on statistical similarities without knowing the original assignment of individuals. The model demonstrates that it can generate images based on the data it has learned, this images are shown in 4.1.



Figure 4.1: Plot of the Faces generated with the GMM.

Another success of the model is evident in anomaly detection:

- The normal face achieves a score of about **-6532413.28**.
- The manipulated (inverted) face scores significantly lower, at about **-33042532.99**.

This massive difference in values of the following picture 4.2 demonstrates that the GMM has successfully learned the underlying structure of human faces. The inverted image is rated by the model as extremely unlikely and is thus reliably detected as an anomaly. This shows that despite the low classification accuracy, the GMM is a powerful tool for evaluating the data.




Figure 4.2: Plot of the 'normal' Faces and the Faces with the anomaly.

5 Conclusion

The analysis of the Olivetti Faces dataset demonstrates the effectiveness of combined machine learning methods. While Q1 introduced structural clustering using K-Means, Q2 was able to significantly improve classification accuracy in the pipeline through a systematic silhouette analysis. In Q3, anomaly detection using GMM and PCA was introduced. The complete source code and many more examples are available on GitHub at https://github.com/luheiss/MachineLearning-DataScience_exercise.

List of Figures

2.1	Plot of the Faces in the first Cluster.	3
3.1	Plot of the Silhouette Score test.	6
4.1	Plot of the Faces generated with the GMM.	9
4.2	Plot of the 'normal' Faces and the Faces with the anomaly.	10

List of Tables

2.1	Distribution of Images per Cluster	3
3.1	Silhouette Scores for different Cluster Sizes (k)	6

References

- [1] scikit-learn developers, “scikit-learn,” 2025. [Online]. Available: https://scikit-learn.org/stable/auto_examples/cluster/plot_kmeans_silhouette_analysis.html#sphx-glr-auto-examples-cluster-plot-kmeans-silhouette-analysis-py
- [2]ageron, “Hands-on machine learning with scikit-learn, keras and tensorflow,” 22. [Online]. Available: https://github.com/ageron/handson-ml2/blob/master/09_unsupervised_learning.ipynb
- [3] P. Daniel T. McGuiness, “Data science ii,” 2025. [Online]. Available: <https://dtmc0945.github.io/L-MCI-BSc-Data-Science-II/DataSciencellLectureBookch5.html#x7-760005.3.1>